



UNIVERSITAT DE BARCELONA



Introduction to HBASE

Advanced Databases

Enric Biosca Trias ebiosca@maia.ub.es

Dept. Matemàtica Aplicada i Anàlisi.

Universitat de Barcelona

Introduct to HBase

- What is HBase?
- Strengths of HBase
- HBase in Production
- Weaknesses of HBase

What is HBase?

- **HBase is a NoSQL database that runs on top of HDFS**
- **HBase is...**
 - Highly available and fault tolerant
 - Very scalable, and can handle high throughput
 - Able to handle massive tables
 - Well suited to sparse rows where the number of columns varies
 - An open-source, Apache project
- **HDFS provides:**
 - Fault tolerance
 - Scalability

- **HBase is based on Google's BigTable**
- **The Google use case is to search the entire Internet**
 - BigTable is at the center of how the company accomplishes this
- **Engineering considerations for searching the Internet:**
 - How do you efficiently download the Web pages?
 - How do you store the entire Internet?
 - How do you index the Web pages?
 - How do you efficiently search the Web pages?



Hbase Principal value?

- **HBase helps solve data access issues where random access is required**
- **HBase scales easily, making it ideal for Big Data storage and processing needs**
- **Columns in an HBase table are defined dynamically, as required**

- **High capacity**

- Massive amounts of data
 - Hundreds of gigabytes, growing to terabytes or even petabytes
- Example: Storing the entire Internet

- **High read and write throughput**

- 1000s/second per node
- Example: Facebook Messages
 - 75 billion operations per day
 - Peak usage of 1.5 million operations per second
 - Runs entirely on HBase

- **Scalable in-memory caching**
 - Adding nodes adds to available cache
- **Large amount of stored data, but queries often access a small subset**
 - Data is cached in memory to speed up queries by reducing disk I/O
- **Data layout**
 - HBase excels at key lookup
 - No penalty for sparse columns

■ Use HBase if...

- You need random write, random read, or both (but not neither)
- Your application performs thousands of operations per second on multiple terabytes of data
- Your access patterns are well-known and relatively simple

■ Don't use HBase if...

- Your application only appends to your dataset, and tends to read the whole thing when processing
- Primary usage is for ad-hoc analytics (non-deterministic access patterns)
- Your data easily fits on one large node

- **Many enterprises are using HBase in production**

- Many use cases are mission critical

- **Examples of companies using HBase:**

- eBay
- Facebook
- FINRA
- Pinterest
- Salesforce
- StumbleUpon
- TrendMicro
- TwiCer
- Yahoo!
- More complete list at

[http://wiki.apache.org/hadoop/Hbase/ PoweredBy](http://wiki.apache.org/hadoop/Hbase/PoweredBy)

Messaging

- Facebook Messages
- Telco SMS/MMS services
- Feeds like Tumblr, Pinterest

Simple Entities

- Geolocation data
- Search index building

Graph Data

- Sessionization
 - Financial transactions
 - Click streams
 - Network traffic

Metrics

- Campaign impressions
- Click counts
- Sensor Data

Messaging

- Facebook Messages
- Telco SMS/MMS services
- Feeds like Tumblr, Pinterest

Characteristics

- Real time random writes
- Reading of top N entries

Simple Entities

- Geolocation data
- Search index building

Characteristics

- Batch or real time, random writes
- Real-time, random reads

Graph Data

- Sessionization
 - Financial transactions
 - Click streams
 - Network traffic

Characteristics

- Batch/real time, random reads/writes

Metrics

- Campaign impressions
- Click counts
- Sensor Data

Characteristics

- Frequently updated metrics

- **HBase does not provide certain popular RDBMS features**
 - Integrated support for SQL
 - External projects such as Hive, Impala, and Phoenix provide various ways of using SQL to access HBase tables
 - Support for transactions
 - Multiple indexes on a table
 - ACID compliance



UNIVERSITAT DE BARCELONA



Different Design Approach in HBase

- **Designing an HBase application requires developers to engineer the system using a data-centric approach**
 - This is a less familiar approach; relationship-centric is more common
 - We will cover this in detail later in the course

- There are many technological challenges when dealing with Big Data
- Hadoop is a distributed system aimed at managing Big Data
- Hadoop's core components are HDFS and MapReduce
- HBase is a NoSQL database that runs on top of HDFS
- HBase is not a traditional RDBMS
- HBase requires a data-centric design approach
- HBase allows for large tables and high throughput



UNIVERSITAT DE BARCELONA



HBASE TABLES

HBase Tables

- **HBase Concepts**
- HBase Table Fundamentals
- Thinking About Table Design
- Hands-On Exercise: HBase Data Import
- Essential Points



■ Definitions

- *Node*
 - A single computer
- *Cluster*
 - A group of nodes connected and coordinated by certain nodes to perform tasks
- *Master Node*
 - A node performing coordination tasks
- *Worker Node*
 - A node performing tasks assigned to it by a master node
- *Daemon*
 - A process or program that runs in the background



Fundamental HBase Concepts

- **HBase stores data in tables**
 - Similar to RDBMS tables but with some important differences
- **Table data is stored on the Hadoop Distributed File System (HDFS)**
 - Data is split into HDFS blocks and stored on multiple nodes in the cluster



What is a Table?

- HBase tables are comprised of rows, columns, and column families
- Every row has a *row key* for fast lookup
- Columns hold the data for the table
- Each column belongs to a particular *column family*
- A table has one or more column families

- **HBase is essentially a distributed, sorted map**
- **Distributed:** HBase is designed to use multiple machines to store and serve table data
- **Sorted Map:** HBase stores table data as a map, and guarantees that adjacent keys will be stored next to each other on disk

- **In our examples, we will use a table that holds**
 - User contact information
 - Profile photos
 - User sign-in information such as username and password
 - Settings or preferences for multiple applications
- **The table will be designed to provide access to the data based on the username**

- **Not every field will have a value**
 - This might happen if the application that stores data to a given field has not run
 - Alternatively, the user may have elected not to provide all the information
- **For now we will focus on the contact information and profile photo**
 - Contact information
 - First Name
 - Last Name
 - Profile photo
 - Image that the user uploads



- **Tables are comprised of rows, columns, and column families**
- **Every row has a *row key***
 - A row key is analogous to a primary key in a traditional RDBMS
 - Rows are stored sorted by row key to enable speedy retrieval of data

- **Columns hold the data for the table**
 - Columns can be created on the fly
 - A column exists for a particular row only if the row has data in that column
 - The table's *cells* (row-column intersection) are arbitrary arrays of bytes
- **Each column in an HBase table belongs to a particular *column family***
 - A collection of columns
- **A table has one or more column families**
 - Created as part of the table definition



HBase Columns and Column Families

- **All columns belonging to the same column family have the same prefix**
 - e.g., `contactinfo:fname` and `contactinfo:lname`
 - The “:” delimits the column family from the qualifier (column name)
- **Tuning and storage settings can be specified for each column family**
 - For example, the number of versions of each cell which will be stored
 - More on this later
- **A column family can have any number of columns**
 - Columns within a family are sorted and stored together

HBase Tables: A Conceptual View

- A column is referenced using its column family and column name (or *qualifier*)
- Separate column families are useful for
 - Data that is not frequently accessed together
 - Data that uses different column family options, e.g., compression

Column Families		contactinfo		profilephoto
Column Names	Row Key	fname	lname	image
Rows	jdupont	Jean	Dupont	
	jsmith	John	Smith	<smith.jpg>
	mrossi	Mario	Rossi	<mario.jpg>



Storing Data in Tables

- **Data in HBase tables is stored as byte arrays**
 - Anything that can be converted to an array of bytes can be stored
 - Strings, numbers, complex objects, images, etc.
- **Cell size**
 - Practical limit on the size of values
 - In general, cell size should not consistently be above 10MB

Table Storage On Disk: Basics

- Data is physically stored on disk on a per-column family basis
- Empty cells are not stored

File

contactinfo		
Row Key	fname	lname
jdupont	Jean	Dupont
jsmith	John	Smith
mrossi	Mario	Rossi

File

profilephoto	
Row Key	image
jsmith	<smith.jpg>
mrossi	<mario.jpg>

■ HBase table features

- Row key + column + timestamp --> value
- Row key and value are just bytes
- Can store anything that can be serialized into a byte array

Sorted
by
Row Key
and
Column

Row Key	Column	Timestamp	Cell Value
jdupont	contactinfo:fname	1273746289103	Jean
jdupont	contactinfo:lname	1273878447049	Dupont
jsmith	contactinfo:fname	1273516197868	John
jsmith	contactinfo:lname	1273871824184	Smith
mrossi	contactinfo:fname	1273616297446	Mario
mrossi	contactinfo:lname	1273971921442	Rossi

- **All rows in HBase are identified by a row key**
 - This is like the primary key in a relational database
- **Get/Scan retrieves data**
 - A **Get** retrieves a single row using the row key
 - A **Scan** retrieves all rows
 - A **Scan** can be constrained to retrieve all rows between a start row key and an end row key
- **Put inserts data**
 - A **Put** adds a new row identified by a row key
 - Multiple **Put** calls can be run to insert multiple rows with different row keys

- **Delete** marks data as having been deleted
 - A **Delete** removes the row identified by a row key
 - The data is not removed from HDFS during the call but is marked for deletion
 - Physical deletion from HDFS happens later
- **Increment** allows atomic counters
 - Cells containing a value stored as a 64-bit integer (a **long**)
 - **Increment** allows the value to be initially set, or incremented if it already has a value
 - Atomicity allows for concurrent access from multiple clients without fear of corruption by a write from another process



UNIVERSITAT DE BARCELONA

Row Key is the Only Indexed Column

- **RDBMSs can have as many index columns as required**
- **In HBase, we have just one indexed column – the row key**
- **Significant effort goes into the row key planning for HBase tables**
 - We rely on the row key to provide quick access to data for all applications that use a given table

Features Comparison: RDBMS vs. HBase

	RDBMS	HBase
Data layout	Row- or column-oriented	Column family-oriented
Transactions	Yes	Single row only
Query language	SQL	get/put/scan
Security	Authentication/ Authorization	Access control at per-cell level, also at cluster, table, or row level
Indexes	Yes	Row key only
Max data size	TBs	PB+
Read/write throughput limits	1000s queries/second	Millions of queries/second



Replacing RDBMSs with HBase

- **Replacing an RDBMS-based applicaCon with HBase requires significant re-architecting**
- **Some major differences**
 - Data layout
 - Data access

■ Joins

- In relational databases, one would typically normalize tables and use joins to retrieve data
- HBase does not support explicit joins
 - Instead, a lookup by row key joins data from column families

■ Scaling

- Relational tables can scale through partitioning or sharding data
- HBase automatically partitions data into smaller pieces

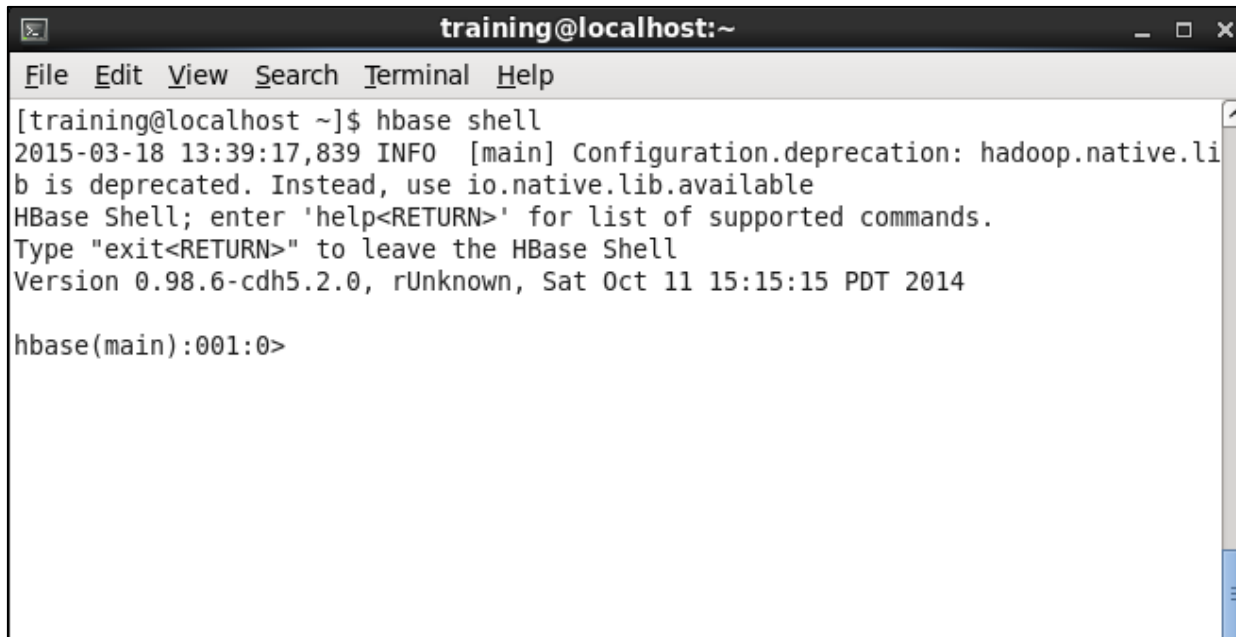
- **Steps for designing an RDBMS schema**
 - Determine all the types of data to be stored
 - Relationship-centric: Determine relationships between data elements
 - Create tables, columns, and foreign keys to maintain relationships
- **Steps for designing an HBase schema**
 - Data-centric: Identify ways in which data will be accessed
 - Identify types of data to be stored
 - Create data layouts and keys

- **The HBase Shell is an interactive shell for sending commands to HBase**
- **HBase shell uses JRuby**
 - Wraps Java client calls in Ruby
 - Allows Ruby syntax to be used for commands
 - Makes parameter usage a little different than most shells
 - Command parameters are single quoted (')



Running the HBase Shell

```
$ hbase shell
```



```
training@localhost:~  
File Edit View Search Terminal Help  
[training@localhost ~]$ hbase shell  
2015-03-18 13:39:17,839 INFO [main] Configuration.deprecation: hadoop.native.lib  
is deprecated. Instead, use io.native.lib.available  
HBase Shell; enter 'help<RETURN>' for list of supported commands.  
Type "exit<RETURN>" to leave the HBase Shell  
Version 0.98.6-cdh5.2.0, rUnknown, Sat Oct 11 15:15:15 PDT 2014  
  
hbase(main):001:0>
```

- **Get help:**

```
hbase> help
```

- **Get HBase status:**

```
hbase> status  
3 servers, 0 dead, 1.3333 average load
```

- **Get version:**

```
hbase> version  
0.98.6-cdh5.2.0, rUnknown, Sat Oct 11 15:15:15 PDT 2014
```

- Shell commands are often followed by parameters

```
hbase> command 'parameter1', 'parameter2'
```

- More advanced parameters use Ruby hashes

- Ruby hash syntax is `{PARAM => 'stringvalue'}`
- The Ruby '`=>`' operator is called a hash rocket

- Commands with advanced parameters look like this:

```
hbase> command 'parameter1', {PARAMETER2 => 'stringvalue',  
PARAMETER3 => intvalue}
```

- Example:

```
hbase> create 'movie', {NAME => 'desc', VERSIONS => 5}
```


- The HBase Shell works in interactive and batch modes
 - Allows a script to be written in JRuby/Ruby and passed into the shell
- Passing scripts to the HBase Shell

```
$ hbase shell pathtorubyscript.rb
```

- Shell scripts can take command line parameters to expand functionality

```
hbase> require 'pathtorubyscript.rb'
```

```
hbase> myRubyFunction 'parameter1'
```

HBase Table Creation

- **Only tables and column families need to be specified at table creation**
 - You must supply names for the table and column family/families
 - Every table must have at least one column family
 - Any optional settings can be changed later
- **Through a new namespace feature, HBase tables can be grouped**
 - Similar to the “database” or “schema” concept in relational database systems
- **Table creation is different from traditional RDBMS table creation**
 - No columns or strict relationships need to be created
 - No constraints or foreign key relationships are specified
 - No database namespace needs to be supplied

- General form:

```
create 'tablename', {NAME => 'colfam' [, options]} [, {...}]
```

- Examples:

```
create 'movie', {NAME => 'desc'}  
  
create 'movie', {NAME => 'desc', VERSIONS => 2}  
create 'movie', {NAME => 'desc'}, {NAME => 'media'}
```

- Shorthand (with default options):

```
create 'movie', 'desc', 'media'
```

Managing HBase Namespaces

- HBase provides the capability to define and manage namespaces
- Multiple tables can belong to a single namespace
 - Define the table's namespace when the table is created

- Namespaces can be created, dropped, or altered:

```
create_namespace 'namespaceName'

drop_namespace 'namespaceName'

alter_namespace 'namespaceName', {METHOD => 'set',
'PROPERTY_NAME' => 'PROPERTY_VALUE'}
```

- There are two predefined special namespaces
 - **hbase** – A system namespace that holds HBase internal tables
 - **default** – A namespace for tables with no explicit namespace defined

Creating Tables in Namespaces

- General form:

```
create 'namespace:tablename', {NAME => 'colfam' [, options]}  
[, {...}]
```

- Examples:

```
create_namespace 'entertainment'  
  
create 'entertainment:movie', {NAME => 'desc'}
```



Listing and Describing Tables

- **List all the tables in HBase**

```
hbase> list
```

- **Give extended details about a table**

- Provides all column families in a table, their properties, and values

```
hbase> describe 'movie'
```

- **Disabling a table puts it in a maintenance state**
 - Allows various maintenance commands to be run
 - Prevents all client access
 - May take up to several minutes for a table to disable

```
hbase> disable 'movie'
```

- **Take the table out of maintenance state**

```
hbase> enable 'movie'
```

- The **drop** command removes the table from HBase and deletes all its files in HDFS
 - The table must be disabled first

```
hbase> disable 'movie'
```

```
hbase> drop 'movie'
```

- The **truncate** command deletes every row in the table
 - Table and column family schema are unaffected
 - It is not necessary to manually disable the table first

```
hbase> truncate 'movie'
```


- Tables can be changed after creation
 - The entire table can be modified
 - Column families can be added, modified, or deleted
- For some operations, tables must be disabled while the changes are being applied
 - Re-enable the table after changes are complete
 - As of HBase 0.92, schema changes such as column family changes do not require the table to be disabled
 - By default, online schema changes are enabled
 - – The `hbase.online.schema.update.enable` property is set to `true`



Adding, Deleting, and Modifying Column Families

■ `hbase> alter 'movie', NAME => 'media'`

■ `hbase> alter 'movie', NAME => 'media', METHOD => 'delete'`

■ `hbase> alter 'movie', NAME => 'desc', VERSIONS => 5`



Altering Column Families Asynchronously

- Use the shell command `alter_async` to alter a column family

-
-

```
hbase> alter_async 'movie', NAME => 'desc', VERSIONS => 6
```

-