



UNIVERSITAT DE BARCELONA



NoSQL Databases – Graph Databases & Neo4j

Advanced Databases


Enric Biosca Trias ebiosca@maia.ub.es

Dept. Matemàtica Aplicada i Anàlisi.

Universitat de Barcelona




Trend 1: Data Size






Trend 2: Connectedness







Side note: RDBMS performance



Key-Value




Graph DB



BigTable

				1	
1				1	
	1		1		
	1	1			
				1	
	1			1	
	1			1	
		1		1	
				1	


Document



- Strengths
 - Powerful data model
 - Fast
 - For connected data, can be many orders of magnitude faster than RDBMS
- Weaknesses:
 - Sharding
 - Though they *can* scale reasonably well
 - And for some domains you can shard too!

- A graph data structure consists of a finite set of **nodes** (or nodes), together with a set of pairs of these vertices. These pairs are known as **edges**, arcs, or lines for an undirected graph and as arrows, directed edges, directed arcs, or directed lines for a directed graph. The vertices may be part of the graph structure, or may be external entities represented by integer indices or references.

- **Set of vertices**
 - Unique id
 - Edges
 - Labels
 - Properties
- **Set of edges**
 - Unique id
 - 2 vertices
 - Label
 - Collection of properties



- **As Storage or Database**
- Graph databases can be used instead of RDBMS.
- **Focused in relationships between entities** rather than attributes.
- **All data can be stored in a single graph.** No need to split it across multiple tables.
- Using pointers it **avoids doing multiple joins.**

Use cases

SQL-like graph queries

- **Operational false positive detection.**
- **speed. NIF-NIF person connections**

Native graph storage

- **Access management system** (metadata)
- **360° view of your data** (all data in a single graph)
- **Non structured data.** Can evolve across time

- Batch analítics engine
- **Computationally intensive tasks.**
- **Complex queries.**
- **Multiple level deep searches.**
- They benefit from parallel and distributed processing.
- **Graph algorithms are more efficient in a graph-native platform.** In RDBMS those algorithms do not always finish.

Use cases

Client clustering (target group marketing)

Ring detection for fraud analysis

Product recommendation

Node centrality, Page Rank to find the most influencer node in a network

Shortest path for traffic optimization system

- **Discovery data**
- It allows the data scientist to **work with all the data at one place**. No need to split the data across different tables.
- **Free schema**. No fixed relationships.
- **Graph visualization** tools are discovery layer's **core component**.
- **Intuitive and easy to use** for nonscientists.

Use cases

Discovery of implicit or latent relationships in data


- **New business opportunities**
- **Custom fraud detection algorithms**
- **Exploration by non tech-focused users**. Example: Panama Papers

Network and IT operations (maintenance, cyber security)

Database



Analytics





Pros

Very popular
Graph pattern-matching language for queries (CYPHER)
Graphical developer interface

Cons

Non distributed solution
No built-in analytics
Slow data ingestion

Connectivity:



Oracle Labs PGX



Pros	Cons
<ul style="list-style-type: none">Can be distributedAnalytics-focused frameworkBuilt-in parallel graph algorithmsCustom algorithms easily parallelizableSQL-like graph querying languageMultiple languages (Java, groovy, scala, R)	<ul style="list-style-type: none">Expertise required for low-level accessHomogeneous graphRAM limited in single-node option

Connectivity:





Spark GraphX



Pros

Distributed
Use already existing spark clusters
Very popular
Apache License
Built-in algorithms
Multiple languages (scala, python, java,R)

Cons

Slower(it is a general-purpose framework)
Need to find optimized algorithms

Connectivity:



mongoDB



S3



Graph exploration

Discovery tools

- Notebooks (Zeppelin, jupyter)

Ad-hoc applications

- Java

Through APIs

Visualization

Desktop applications

- Already developed, only need to be supplied with data
- Custom functionality limited to plugins

Web libraries

- Compatible with any modern browser
- Must be developed specifically for our needs

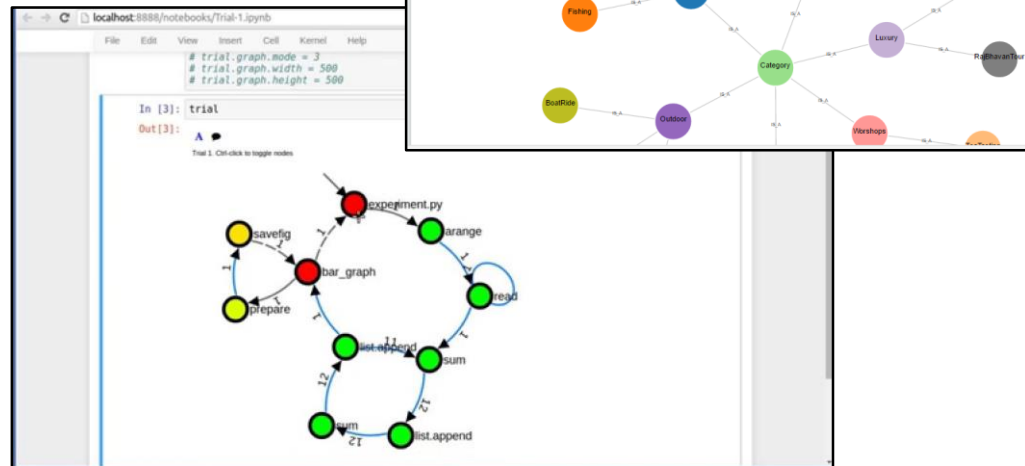
• Discovery t



- **Web-based notebooks enables interactive analytics.**
- **Open-source**

Common features:

- Pluggable visualization frameworks
- Interactive and collaborative
- Allows any language / data-processing backend.
- Supports many interpreters (Java, Scala, Python...)
- REST API


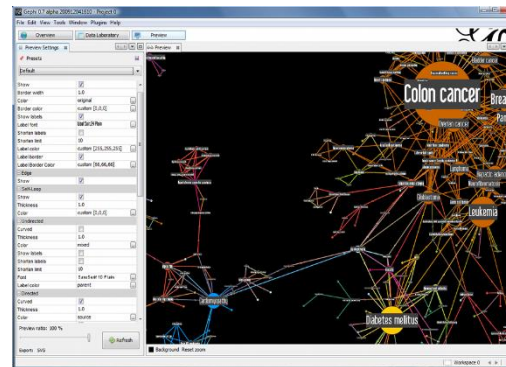


- Desktop applications
- *Gephi*
- Network analysis and visualization software
- Open-source



Common features:

- Customizable visual format
- Built-in layout algorithms
- Statistics and metrics computation
- Node/edge filtering
- Support plugins



- **Web libraries**
- **There are numerous web libraries capable of graph representation and manipulation.**
- **Customized interaction (such as layout algorithms or specific graph operations) must often be developed.**
- **Most notorious options are D3.js and Linkurious**



sigma.js



plotly.js

arbor.js

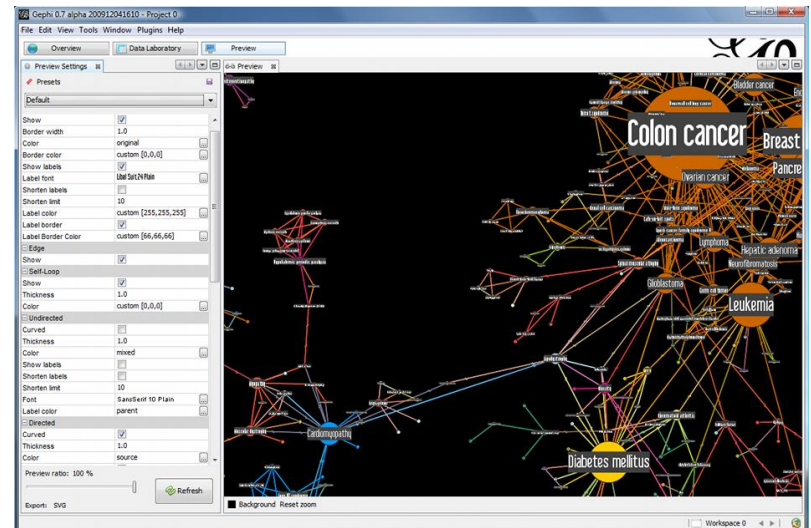
- Desktop applications



- Network analysis and visualization software
- Open-source
- Developed by *The Gephi Consortium*, a non-profit corporation.

Main features:

- Real-time visualization (up to 1M nodes)
- Built-in layout algorithms
- Statistics and metrics computation
- Dynamic node/edge filtering
- Supports plugins




• Desktop applications



- **Molecular network analysis software**
- **Open-source**
- **Can automatically deploy web-based visualizations of graphs loaded in the desktop application**

Main features:

- Supports many standard network file formats
- Built-in layout algorithms
- Node/edge filtering
- Supports plugins
- Available as both the desktop application (Cytoscape) and web libraries (Cytoscape.js)



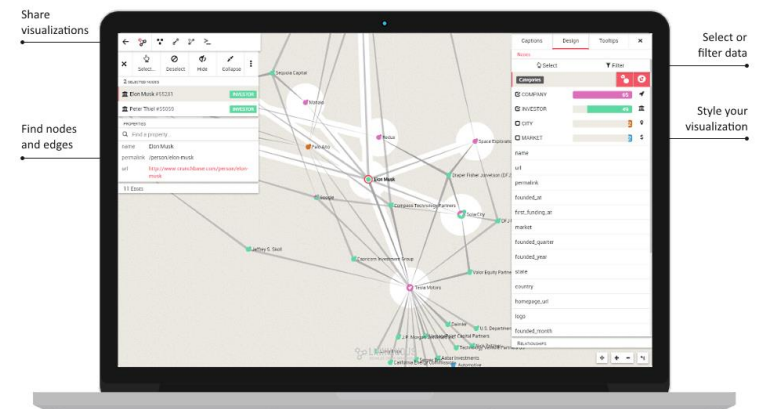


• Web libraries

- **Commercial JavaScript library based on Sigma.js**
- **Focuses on providing a business intelligence approach to graphs**
- **Provides business support and integrations for companies**

Main features:

- Rendering for up to 20K nodes
- Supports data import/export operations with numerous formats
- Plugins available
- Built-in search engine
- Cypher language support to query the visualized graph
- Built-in graph data editor




- Web libraries

Data-Driven Documents

- JavaScript library to create graphic representations of data
- Not a graph-specific library, but can display graphs with great user interaction and dynamic modifications.

Main features:

- Excels at graphic customization
- Supports animated transitions
- Large community support



- Experiment:
 - ~1k persons
 - Average 50 friends per person
 - `pathExists(a, b)` limited to depth 4
 - Caches warm to eliminate disk IO

	# persons	query time
Relational database	1000	2000ms

- Experiment:
 - ~1k persons
 - Average 50 friends per person
 - `pathExists(a, b)` limited to depth 4
 - Caches warm to eliminate disk IO


	# persons	query time
Relational database	1000	2000ms
Neo4j	1000	2ms

- Experiment:
 - ~1k persons
 - Average 50 friends per person
 - `pathExists(a, b)` limited to depth 4
 - Caches warm to eliminate disk IO

	# persons	query time
Relational database	1000	2000ms
Neo4j	1000	2ms
Neo4j	1000000	2ms




Property Graph Model



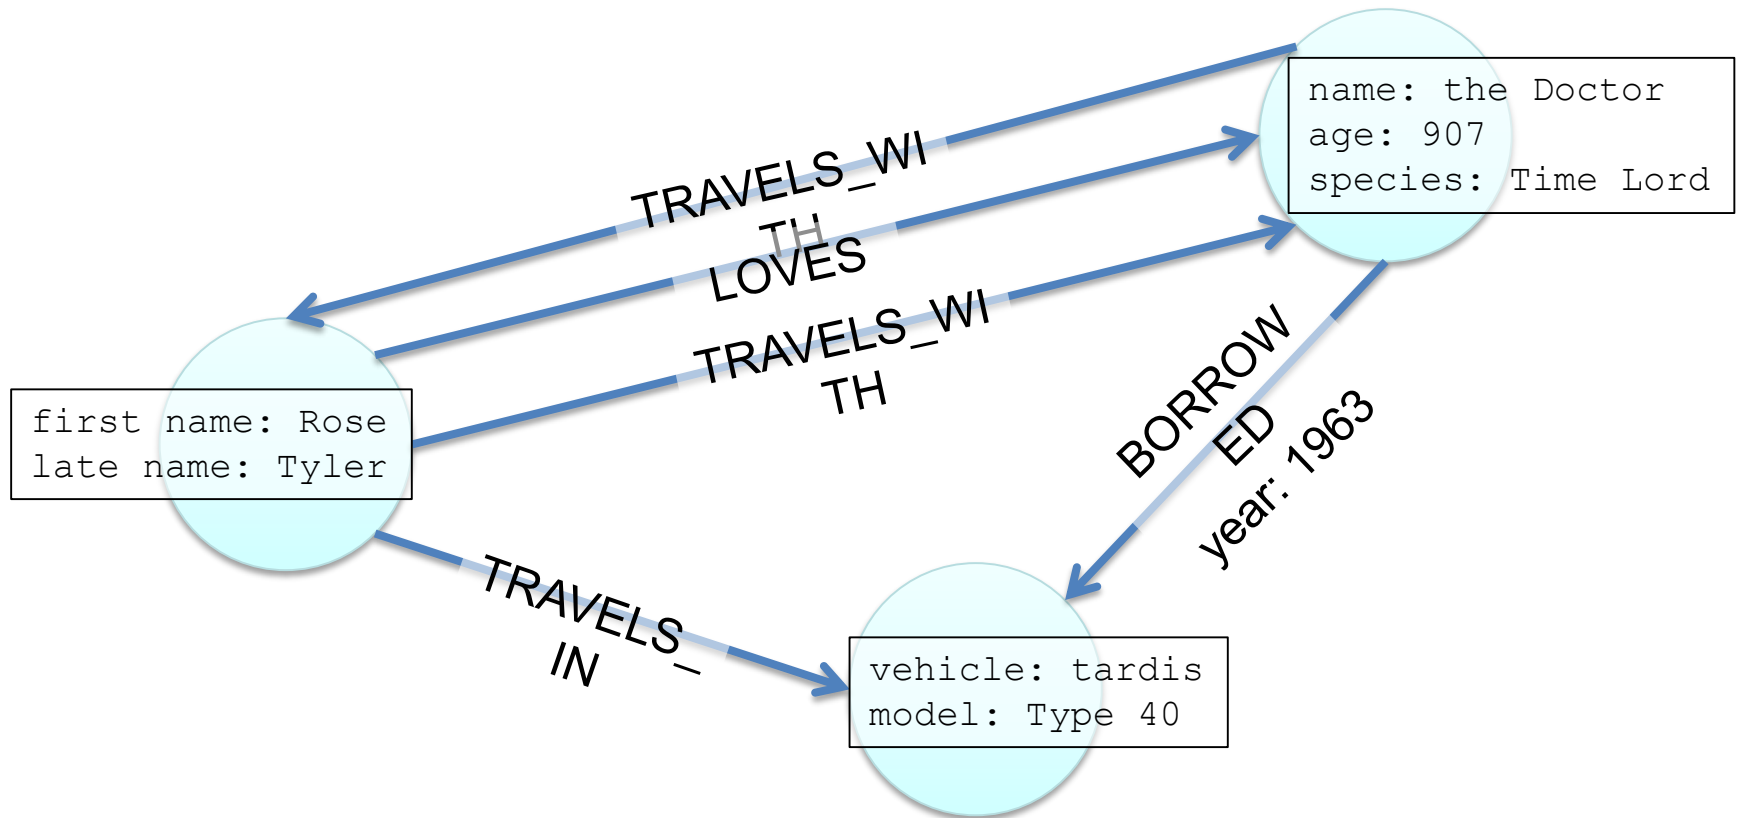


Property Graph Model



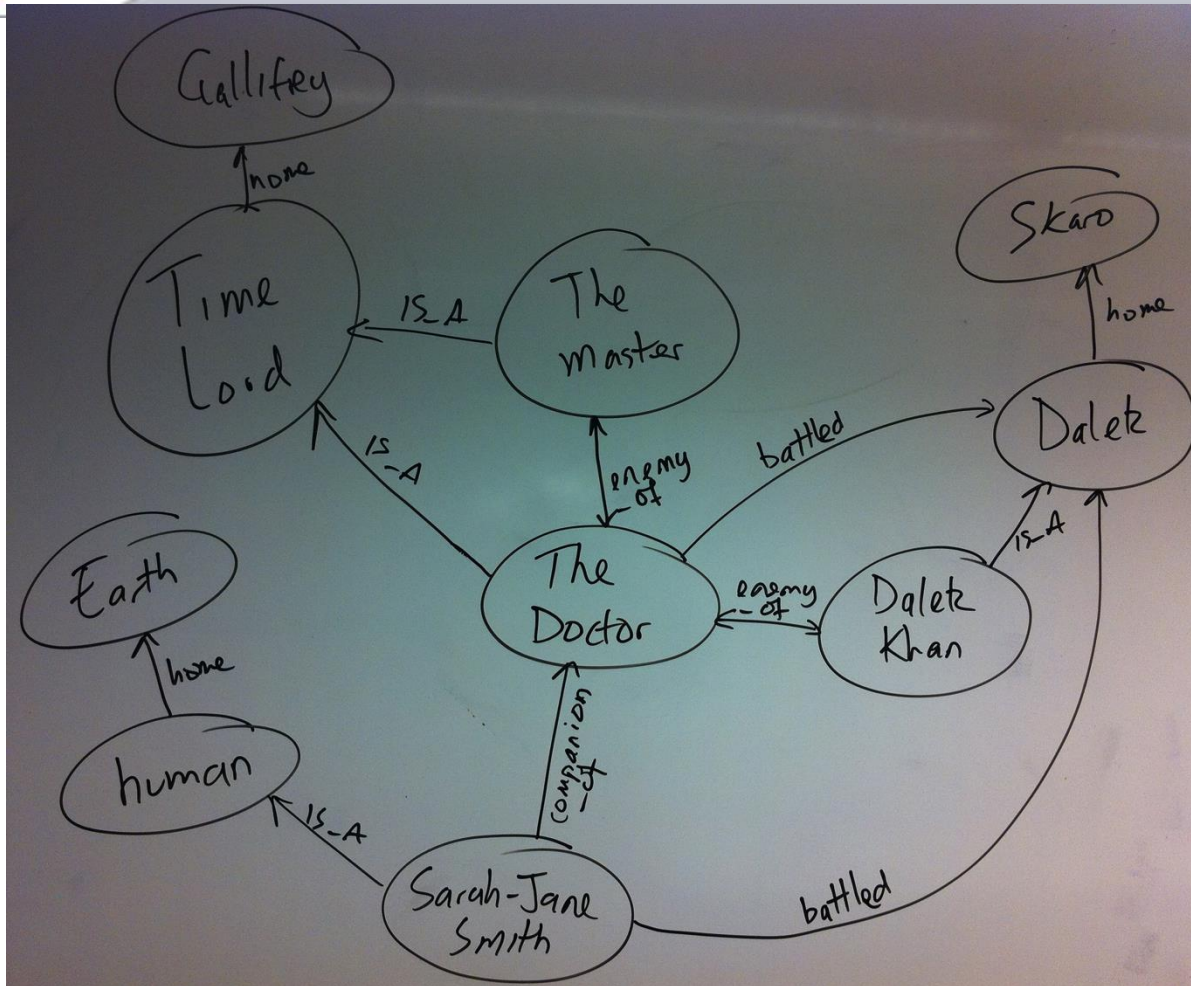


Property Graph Model





Graphs are very whiteboard-friendly



- Graph databases don't excuse you from design
 - Any more than dynamically typed languages excuse you from design
- Good design still requires effort
- But less difficult than RDBMS because you don't need to normalise
 - And then de-normalise!



UNIVERSITAT DE BARCELONA



neo4j

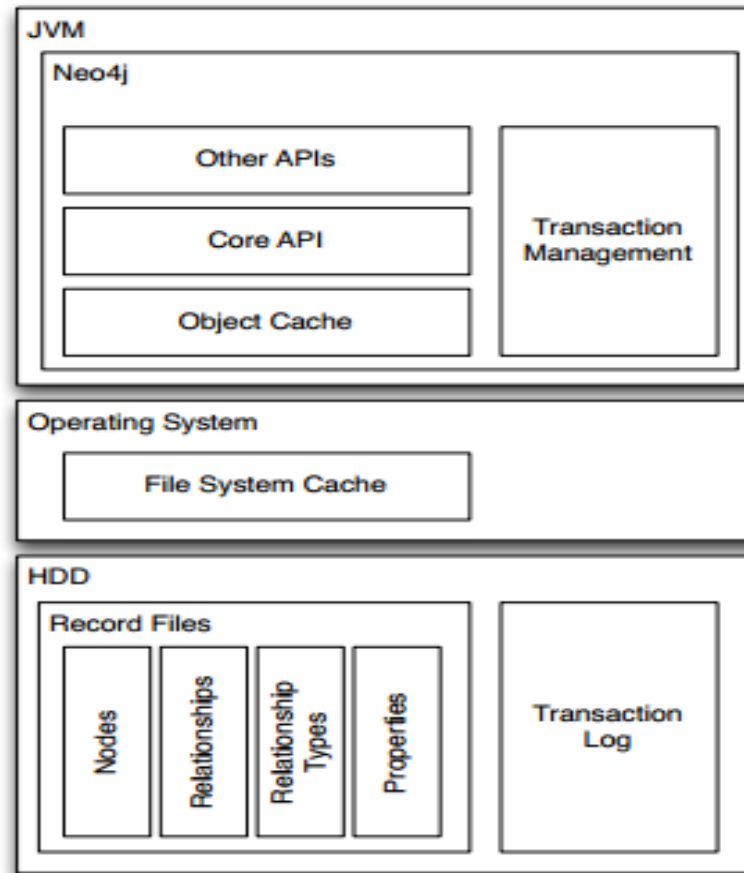
- It's is a Graph Database
- Embeddable and server
- Full ACID transactions
 - don't mess around with durability, ever.
- Schema free, bottom-up data model design



- Neo4j is stable
 - In 24/7 operation since 2003
- Neo4j is under active development
- High performance graph operations
 - Traverses 1,000,000+ relationships / second on commodity hardware



Neo4j Logical Architecture





So how do we query it?

- Through the Java APIs
 - JVM languages have bindings to the same APIs
 - JRuby, Jython, Clojure, Scala...
- Managing nodes and relationships
- Indexing
- Traversing
- Path finding
- Pattern matching

- Deals with graphs in terms of their fundamentals:
 - Nodes
 - Properties
 - KV Pairs
 - Relationships
 - Start node
 - End node
 - Properties
 - KV Pairs

```
GraphDatabaseService db = new
    EmbeddedGraphDatabase("/tmp/neo");
Transaction tx = db.beginTx();
try {
    Node theDoctor = db.createNode();
    theDoctor.setProperty("character", "the
Doctor");
    tx.success();
} finally {
    tx.finish();
}
```

```
Transaction tx = db.beginTx();
try {
    Node theDoctor = db.createNode();
    theDoctor.setProperty("character", "The Doctor");

    Node susan = db.createNode();
    susan.setProperty("firstname", "Susan");
    susan.setProperty("lastname", "Campbell");

    susan.createRelationshipTo(theDoctor,
        DynamicRelationshipType.withName("COMPANION_OF"));

    tx.success();
} finally {
    tx.finish();
}
```

- Graphs are their own indexes!
- But sometimes we want short-cuts to well-known nodes
- Can do this in our own code
 - Just keep a reference to any interesting nodes
- Indexes offer more flexibility in what constitutes an “interesting node”



Indexes trade read performance for write
cost

Just like any database, even RDBMS

**Don't index every
node!
(or relationship)**

- The default index implementation for Neo4j
 - Default implementation for `IndexManager`
- Supports many indexes per database
- Each index supports nodes *or* relationships
- Supports exact and regex-based matching
- Supports scoring
 - Number of hits in the index for a given item
 - Great for recommendations!

- Indexes are typically used only to provide starting points
- Then the heavy work is done by traversing the graph
- Can happily mix index operations with graph operations to great effect

Core

- Basic (nodes, relationships)
- Fast
- Imperative
- Flexible
 - Can easily intermix mutating operations

Traversers

- Expressive
- Fast
- Declarative (mostly)
- Opinionated



- Neo4j has declarative traversal frameworks
 - What, not how
- There are two of these
 - And development is active
 - No “one framework to rule them all” yet

- Mature
- Designed for the 80% case
- In the **`org.neo4j.graphdb`** package

```
Node daleks = ...
Traverser t = daleks.traverse(
    Order.DEPTH_FIRST,
    StopEvaluator.DEPTH_ONE,
    ReturnableEvaluator.ALL_BUT_START_NODE,
    DoctorWhoRelations.ENEMY_OF,
    Direction.OUTGOING);
```



Custom ReturnableEvaluator

```
Traverser t = theDoctor.traverse(Order.DEPTH_FIRST,  
    StopEvaluator.DEPTH_ONE,  
    new ReturnableEvaluator() {  
        public boolean isReturnableNode(TraversalPosition pos)  
        {  
            return pos.currentNode().hasProperty("actor");  
        }  
    },  
    DoctorWhoRelations.PLAYED, Direction.INCOMING);
```

- Newer (obviously!)
- Designed for the 95% use case
- In the `org.neo4j.graphdb.traversal` package
- http://wiki.neo4j.org/content/Traversal_Framework

```
Traverser traverser = Traversal.description()
    .relationships(DoctorWhoRelationships$.ENEMY_OF, Direction.OUTGOING)
    .depthFirst()
    .uniqueness(Uniqueness.NODE_GLOBAL)
    .evaluator(new Evaluator() {
        public Evaluation evaluate(Path path) {
            // Only include if we're at depth 2, for enemy-of-enemy
            if(path.length() == 2) {
                return Evaluation.INCLUDE_AND_PRUNE;
            } else if(path.length() > 2){
                return Evaluation.EXCLUDE_AND_PRUNE;
            } else {
                return Evaluation.EXCLUDE_AND_CONTINUE;
            }
        }
    })
    .traverse(theMaster);
```


- Declarative graph pattern matching language
 - “SQL for graphs”
 - Tabular results
- Cypher is evolving steadily
 - Syntax changes between releases
- Supports queries
 - Including aggregation, ordering and limits
 - Mutating operations in product roadmap

- The top 5 most frequently appearing companions:

```
start doctor=node:characters(name = 'Doctor')
match (doctor)<-[:COMPANION_OF]-(companion)
      -[:APPEARED_IN]->(episode)
return companion.name, count(episode)
order by count(episode) desc
limit 5
```

Start node
from index

Subgraph
pattern

Accumulates
rows by
episode

Limit returned
rows



+-----+		
companion.name	count (episode)	
+-----+		
Rose Tyler	30	
Sarah Jane Smith	22	
Jamie McCrimmon	21	
Amy Pond	21	
Tegan Jovanka	20	
+-----+		
5 rows, 49 ms		
+-----+		

```
ExecutionEngine engine = new ExecutionEngine(database);

String cql = "start doctor=node:characters(name='Doctor') "
            + " match (doctor)<-[:COMPANION_OF]-"
            + "(companion) "
            + "-[:APPEARED_IN]->(episode) "
            + " return companion.name, count(episode) "
            + " order by count(episode) desc limit"
            + " 5";

ExecutionResult result = engine.execute(cql);
```

```
ExecutionEngine engine = new ExecutionEngine(database);

String cql = "start doctor=node:characters(name='Doctor') "
            + " match (doctor)<-[:COMPANION_OF]-"
            + "(companion) "
            + "-[:APPEARED_IN]->(episode) "
            + " return companion.name, count(episode) "
            + " order by count(episode) desc limit"
            + " 5";

ExecutionResult result = engine.execute(cql);
```

Top tip:

`ExecutionResult.dumpToString()`
is your best friend

- Aggregation:

COUNT, SUM, AVG, MAX, MIN, COLLECT

- Where clauses:

```
start doctor=node:characters(name = 'Doctor')
match (doctor)<-[:PLAYED]-(actor)-[:APPEARED_IN]->(episode)
where actor.actor = 'Tom Baker'
      and episode.title =~ /.*Dalek.*/
return episode.title
```

- Ordering:

```
order by <property>
order by <property> desc
```



Cypher Query



```
start daleks=node:species(species='Dalek')
match daleks-[:APPEARED_IN]->episode<-
[:USED_IN]-
      ()<-[:MEMBER_OF]-()-[:COMPOSED_OF]->
      part-[:ORIGINAL_PROP]->originalprop
return originalprop.name, part.type,
count(episode)
order by count(episode) desc
limit 1
```



Index Lookup



```
start daleks=node:species(species='Dalek')
match daleks-[:APPEARED_IN]->episode<-
[:USED_IN]-
      ()<-[:MEMBER_OF]-()-[:COMPOSED_OF]->
      part-[:ORIGINAL_PROP]->originalprop
return originalprop.name, part.type,
count(episode)
order by count(episode) desc
limit 1
```


Match Nodes & Relationships



```
start daleks=node:species(species='Dalek')
match daleks-[:APPEARED_IN]->episode<-
[:USED_IN]-
      ()<-[:MEMBER_OF]-()-[:COMPOSED_OF]->
      part-[:ORIGINAL_PROP]->originalprop
return originalprop.name, part.type,
count(episode)
order by count(episode) desc
limit 1
```

Return Values



```

start daleks=node:species(species='Dalek')
match daleks-[:APPEARED_IN]->episode<-
[:USED_IN]-
      ()<-[:MEMBER_OF]-()-[:COMPOSED_OF]->
      part-[:ORIGINAL_PROP]->originalprop
return originalprop.name, part.type,
count(episode)
order by count(episode) desc
limit 1

```

- Payback time for Algorithms 101
- Graph algos are a higher level of abstraction
 - You do less work!
- The database comes with a set of useful algorithms built-in
 - Time to get some payback from Algorithms 101



- The Doctor and the Master been around for a while
- But what's the key feature of their relationship?
 - They're both timelords, they both come from Gallifrey, they've fought

Try the Shortest Path!

What's the most direct path between the Doctor and the Master?

```
Node theMaster = ...
```

```
Node theDoctor = ...
```

```
int maxDepth = 5;
```

```
PathFinder<Path> shortestPathFinder =  
    GraphAlgoFactory.shortestPath(  
        Traversal.expanderForAllTypes(),  
        maxDepth);
```

```
Path shortestPath =  
    shortestPathFinder.findSinglePath(theDoctor, theMaster);
```



Path finder code

algo

```
Node rose = ...
```

```
Node daleks = ...
```

```
PathFinder<Path> pathFinder = GraphAlgoFactory.pathsWithLength(  
    Traversal.expanderForTypes(  
        DoctorWhoRelationships.APPEARED  
        Direction.BOTH), 2);
```

constraints

fixed path length

```
Iterable<Path> paths = pathFinder.findAllPaths(rose, daleks);
```

- It's super-powerful for looking for patterns in a data set
 - E.g. retail analytics
- Higher-level abstraction than raw traversers
 - You do less work!

Setting up and matching a pattern

[illegible]

- The only way to access the server today
 - Binary protocol part of the product roadmap
- JSON is the default format
 - Remember to include these headers:
 - Accept:application/json
 - Content-Type:application/json

```
curl http://localhost:7474/db/data/node/1
```

```
{
  "outgoing_relationships" : "http://localhost:7474/db/data/node/1/relationships/out",
  "data" : {
    "character" : "Doctor"
  },
  "traverse" : "http://localhost:7474/db/data/node/1/traverse/{returnType}",
  "all_typed_relationships" : "http://localhost:7474/db/data/node/1/relationships/all/{-list|&|types}",
  "property" : "http://localhost:7474/db/data/node/1/properties/{key}",
  "self" : "http://localhost:7474/db/data/node/1",
  "properties" : "http://localhost:7474/db/data/node/1/properties",
  "outgoing_typed_relationships" : "http://localhost:7474/db/data/node/1/relationships/out/{-list|&|types}",
  "incoming_relationships" : "http://localhost:7474/db/data/node/1/relationships/in",
  "extensions" : {
  },
  "create_relationship" : "http://localhost:7474/db/data/node/1/relationships",
  "all_relationships" : "http://localhost:7474/db/data/node/1/relationships/all",
  "incoming_typed_relationships" : "http://localhost:7474/db/data/node/1/relationships/in/{-list|&|types}"
}
```